

Statement of Purpose:

My name is Aryasomayajula Ram Bharadwaj. I am writing to express my interest in pursuing a Ph.D. in Artificial Intelligence, with a focus on AI safety. With over six years of professional experience in machine learning (ML) and MLOps, as well as a strong background in software engineering, I have developed a solid foundation in creating scalable ML solutions. My current research interests include developing methods for better interpret and evaluate LLMs to improve trustworthiness and avoid harmful behaviors. I will elaborate on these interests in the following sections.

Professional Background and Experience:

After completing my B.Tech in Electronics and Communications in 2019 at GMR Institute of Technology, India, I joined Mu Sigma Business Solutions, a leading decision sciences firm, where I played a key role in the innovation and development lab. During my tenure, I contributed to various development projects. Notably, I led the development of an LLM-agent platform, which was successfully deployed for internal use in two research labs of major Fortune 500 companies. This platform was designed to assist entry-level and mid-level data analysts with their projects. Additionally, I was instrumental in developing the company's proprietary trading platform and made substantial contributions to the in-house MLOps platform.

Currently, as a Technical Architect at Quantiphi Analytics, I specialize in building LLM-agent-based chatbot solutions for a diverse range of clients. My ability to integrate technical innovation with practical deployments has been recognized through multiple awards, including the *Star Performer of the Team* and various *Impact Awards* throughout my career. These experiences have provided me with a unique perspective on the challenges and opportunities in deploying advanced AI systems in real-world settings.

Independent Research and Contributions:

In addition to my professional responsibilities, I have actively pursued independent research. My recent paper, *Understanding Hidden Computations in Transformer Language Models: Filler Tokens in Chain-of-Thought Reasoning* [1], examines how filler tokens affect the faithfulness of CoT explanations. This work offers insights into improving the interpretability of CoT reasoning and highlights a relatively underexplored approach in interpreting chain of thought reasoning.

My interest in AI alignment and safety has grown significantly over the past few years. I have received recognition through prizes in the **AI Alignment Awards** [2] and the **ELK challenge**, which reflect my strong motivation and ability to undertake independent research in this critical area. These experiences have inspired me to deepen my understanding of AI systems and their impact on society.

Research Interests:

While I am open to exploring a wide range of problems related to improving our understanding and trustworthiness of modern AI systems, my current research interests focus on the following key areas:

- **Faithfulness in Chain-of-Thought Explanations:** Chain-of-thought reasoning is a simple yet highly effective technique for enabling LLMs to perform complex reasoning tasks. CoT explanations also provide a window into understanding the underlying reasoning mechanisms of these models. However, it remains unclear how faithful these explanations truly are. Recent work, such as **"Let's Think Dot by Dot"** [3], **"Language Models Don't Always Say What They Think"** [4], and other studies on measuring faithfulness in CoT reasoning, has shown that we cannot fully rely on these explanations. My recent paper explores the problem of CoT explanations involving filler tokens. I believe there is significant value in deeply understanding and addressing the limitations of CoT explanations.

- **Evaluations for LLM Agents:** The rise of LLM agents, which combine powerful language models with tools to act and plan in real-world environments, underscores the need to evaluate their capabilities and address potential risks posed by their open-ended and high-stakes applications. However, evaluating LLM agents before deployment in production is a complex challenge. It is therefore crucial to develop robust evaluation frameworks for identifying dangerous capabilities such as sandbagging, deception, and self-awareness.

Why a Ph.D. and Future Goals:

I have actively pivoted my career toward becoming an AI safety researcher over the past few years, driven by my belief that this is one of the most pressing and impactful challenges of our time. Through this journey, I have come to appreciate the immense value of expert guidance and structured mentorship, particularly in a field as dynamic and significant as AI safety. I view a Ph.D. program as the ideal platform to refine my research skills and transition from an enthusiastic practitioner to a dedicated researcher. I am especially interested in working with **Assistant Professor Tianyi Zhou** on understanding limitations of current LLMs from a safety perspective or **Assistant Professor Yizheng Chen** on evaluating LLM Agents for harmful behaviors. I am committed to pursuing a Ph.D. at University of Maryland to grow intellectually and make meaningful contributions to cutting-edge advancements in AI safety.

Conclusion:

My professional background, independent research experience, and passion for addressing challenging problems related to AI safety make me a strong candidate for your Ph.D. program. I am excited about the opportunity to join University of Maryland and contribute to its vibrant academic community. Thank you for considering my application.

Regards,

Aryasomayajula Ram Bharadwaj

References

- [1] Aryasomayajula Ram Bharadwaj. (2024). *Understanding Hidden Computations in Chain-of-Thought Reasoning*. <https://arxiv.org/abs/2412.04537>
- [2] AI Alignment Awards Winners. <https://www.alignmentawards.com/winners>
- [3] Jacob Pfau, William Merrill, and Samuel R. Bowman. (2024). *Let’s Think Dot by Dot: Hidden Computation in Transformer Language Models*. <https://arxiv.org/abs/2404.15758>
- [4] Miles Turpin, Julian Michael, Ethan Perez, and Samuel R. Bowman. (2023). *Language Models Don’t Always Say What They Think: Unfaithful Explanations in Chain-of-Thought Prompting*. <https://arxiv.org/abs/2305.04388>
- [5] Paul Christiano, Ajeya Cotra, and Mark Xu. (2021). *Eliciting Latent Knowledge: How to Tell If Your Eyes Deceive You*. Alignment Research Center, December 2021. https://docs.google.com/document/d/1WwsnJQstPq91_Yh-Ch2XRL8H_EpsnjrC1dwZXR37PC8/edit